

Analysis	Cards	Setup	Result	Takeaway
Determinism (fresh-process re-run)	100	Re-run on a restarted server; greedy decoding (temperature 0, fixed seed). Genuine ~5 s per card decodes, not a cache replay.	100% paired agreement (100/100)	Calls reproduce exactly across a restart.
Prompt sensitivity (3 reworded prompts)	100	Three reworded system prompts on the same gallery images, vs the canonical-prompt reference on those images.	Concordance: compact 66.0%, explicit 69.5%, terse 52.0%. Positive rate: compact 84%, explicit 71%, terse 79% (reference 48% positive over 568).	Calls shift with wording; variants over-call.
Render sensitivity (prompt + render path)	100	Same three variants on gallery images, vs the independent render-on-the-fly reference (prompt and render path both differ).	Concordance: compact 38.0%, explicit 42.1%, terse 44.0%.	Render path adds a second large source of disagreement.